

ROB 311-Task 3


Q-Learning

Adriana TAPUS & Chuang YU

adriana.tapus@ensta-paris.fr & chuang.yu@ensta-paris.fr

09-2020

1. Introduction



Pac-Man

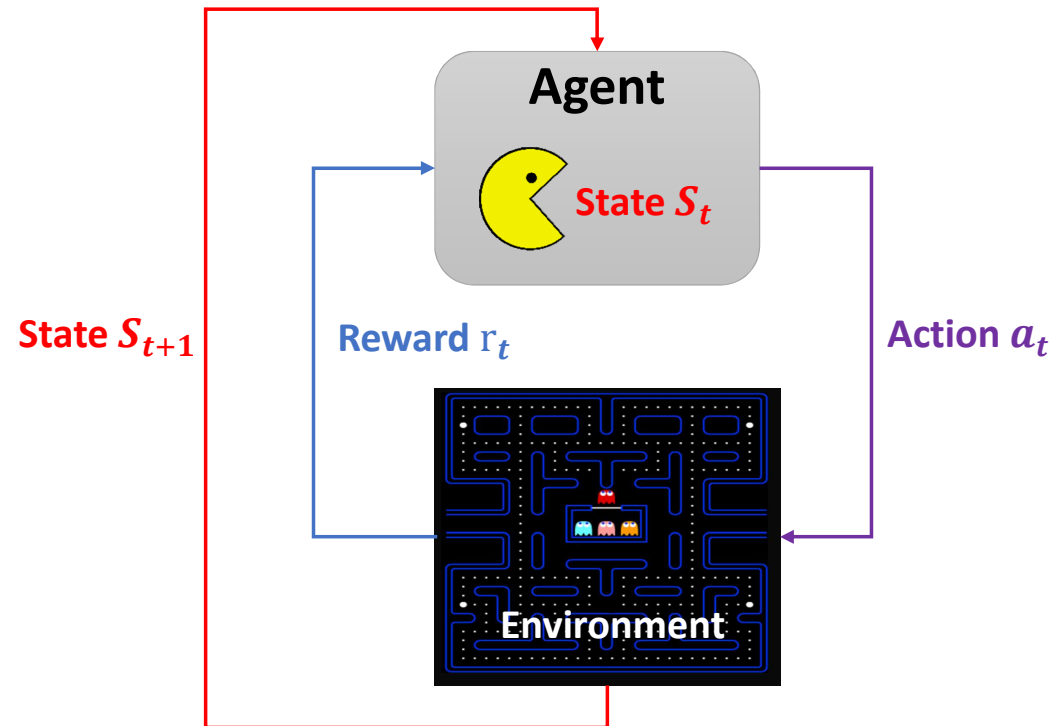
Agent: Player Environment: Game

State S_t : the position at time “ t ”, this frame

Action a_t : $a_t \in \{\text{“Left”, “Right”, “Up”, “Down”}\}$

Reward r_t : at time “ t ”, eat one bean, **immediate**

Reward less (little bean) or more (corner big bean).




Policy π : It is the probability of taking action.

$$\pi(a | s) = P(A = a | S = s)$$

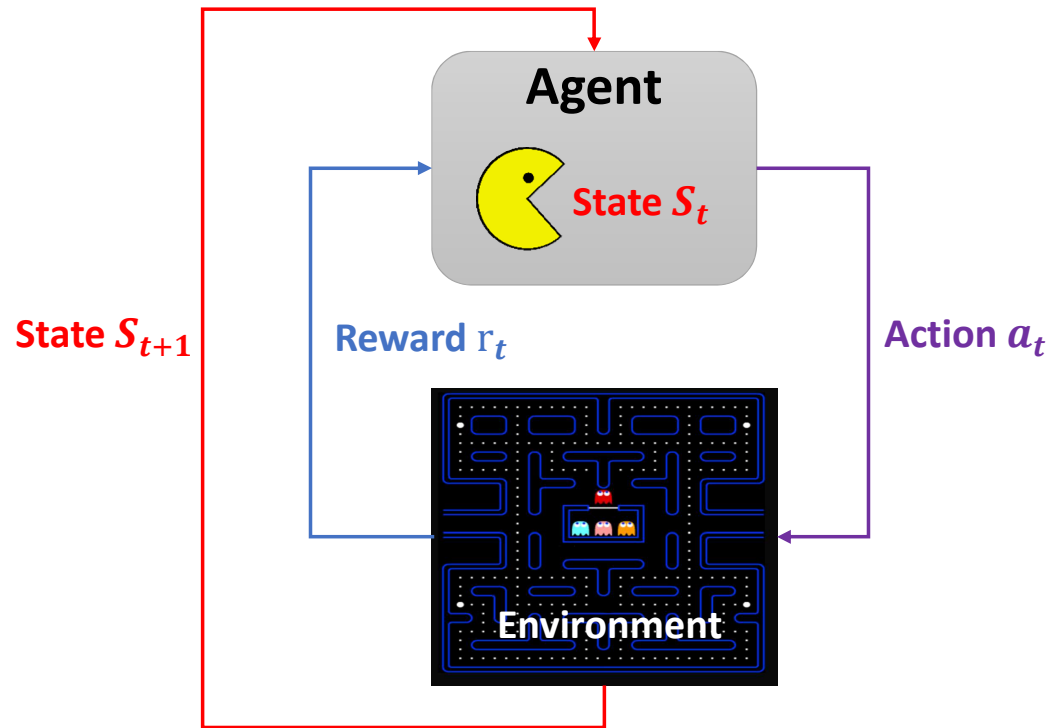
$$\text{e.g. } \pi(\text{Left} | s) = 0.4 \quad \pi(\text{Right} | s) = 0.2$$

$$\pi(\text{Up} | s) = 0.2 \quad \pi(\text{Down} | s) = 0.2$$

1. Introduction



Pac-Man



State Transition: State $S_t \rightarrow$ State S_{t+1}

- PacMan takes **Action a_t** , the **state** changes
- $P(s'|s, a) = P(S_{t+1} = s' | S_t = s, A_t = a)$

1. Introduction

Return : cumulative future reward

$$U_t = R_t + R_{t+1} + R_{t+2} + \dots$$



Discounted Return : discounted cumulative future reward

$$U_t = R_t + \gamma \cdot R_{t+1} + \gamma^2 \cdot R_{t+2} + \dots$$

Action-value function $Q_\pi(s_t, a_t)$

$$Q_\pi(s_t, a_t) = \mathbb{E}[U_t | S_t = s_t, A_t = a_t]$$

State-value function

$$V_\pi(s_t) = \mathbb{E}_A[Q_\pi(s_t, A)] = \sum_a \pi(a|s_t) \cdot Q_\pi(s_t, a)$$

2. Q-Learning

$$V^*(S) = R(s) + \max_a \gamma \sum_{S'} T(S, a, S') V^*(S')$$

Value iteration is a method of computing an **optimal policy** for an MDP (Markov Decision Process) and **its value**.

Value iteration →
$$\begin{cases} Q_{k+1}(s, a) = R(s) + \gamma \sum_{s'} T(S, a, S') * V_k(s') & (1) \\ V_k(s) = \max_a Q_k(s, a) & (2) \end{cases}$$

Replace $V_k(s')$ in Equation (1) with Equation (2)

$$Q_{k+1}(s, a) = \sum_{s'} T(S, a, s') [R(s) + \gamma \max_a Q_k(s', a)]$$

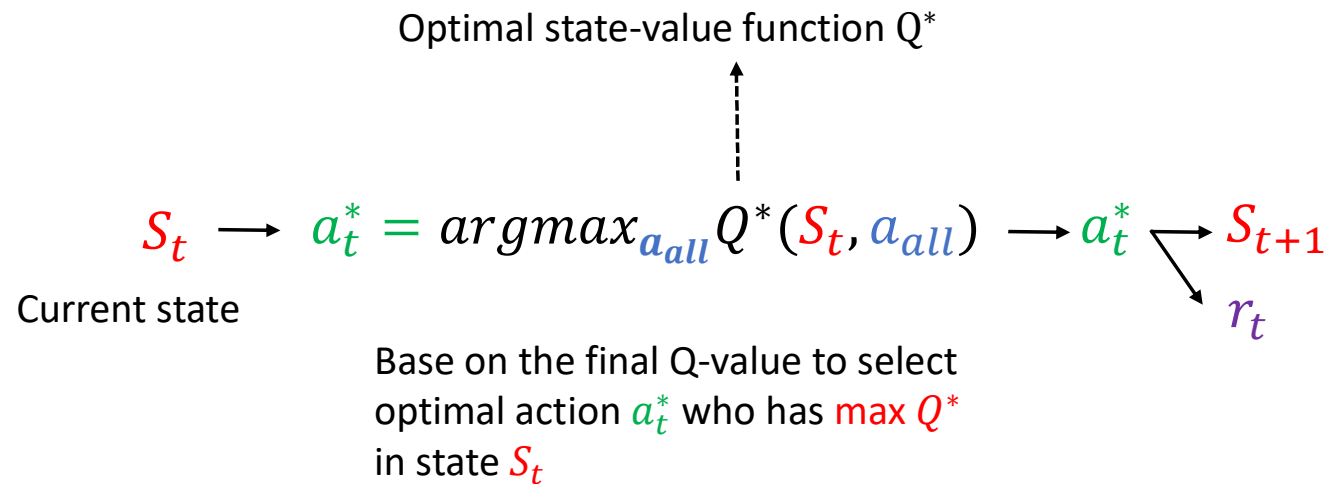
Q-Value update

(Q-learning) →
$$Q(S, A) \leftarrow (1 - \alpha)Q(S, A) + \alpha[R(S, a) + \gamma \max_a Q(S', a)]$$

https://artint.info/html/ArtInt_265.html

2. Q-Learning

Process of Q-Learning playing Pac-Man



3. Task 3 Q-learning



ROB311 – TP3 – Q-Learning and Pac-man

Introduction

Today, you will implement **Q-Learning** in Python. In particular, you will train an AI player of the famous [Pac-Man arcade game](#). You will work on a set of Python files and libraries provided to you by the **UC Berkeley university**, as part of their *CS188 – Intro to AI* course, and you will only have to write the core algorithm of the methods dealing with the Q-Learning and Approximate Q-Learning, based on the equations you have seen in the course.

Files

The **pacman.zip** archive you have downloaded from the ROB311 course page provides you with the files you need in order to implement the Q-Learning and Approximate Q-Learning AI for the Pac-Man game. The files are taken from the [Project 3: Reinforcement Learning](#) page of the *CS188 – Intro to AI* online course. You can read the **Introduction** section of the page if you need extra informations, but we are going to focus only on the files you need in order to implement today's algorithms.

File and classes you have to modify:

- **qlearningAgents.py**: *QLearningAgent, ApproximateQAgent*

AI/Machine Learning related files and classes (you will have to look through):

- **learningAgents.py**: *ReinforcementAgent, ValueEstimationAgent*
- **util.py**: *Counter*
- **game.py**: *Agent*
- **featureExtractors.py**

[Related PPT-Q Learning in CS 188](#)

[Project 3: Reinforcement Learning](#)

3. Task 3 Q-learning



Objectives

Implement Q-Learning by modifying the following *QLearningAgent* methods in the *qLearningAgents.py* file:

- `__init__()`
- `getQValue()`
- `computeValueFromQValues()`
- `computeActionFromQValues()`
- `getAction()`
- `update()`

3. Task 3 Q-learning



How to run and test your code

Navigate using a terminal to your project folder and run the following commands:

```
python pacman.py -p PacmanQAgent -x 2000 -n 2010 -l smallGrid
```

This command will attempt learning from 2000 training episodes and then test the resulting AI agent (player) on 10 games. You will be able to see the AI playing the 10 test games. The learning is done on the *smallGrid* map and the results of the 10 test games will be displayed in the terminal. The win rate on the 10 test game is supposed to be very high, ideally 100%.

```
python pacman.py -p PacmanQAgent -n 10 -l smallGrid -a numTraining=10
```

This command will show you what happens during the training process for 10 games.

If you want to test your code in other scenarios, use the following command to understand what each of the command line parameters does:

```
python pacman.py --help
```



Use Layout "SmallGrid"

Rule

- 2 persons in one group
- Deadline: Before Monday

Submit:

- Report paper + code
- Github or ENSTA gitlab

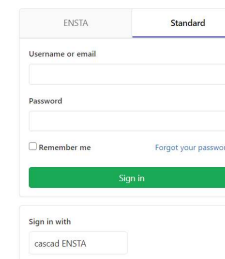


or

Bienvenue sur le serveur GitLab de DaTA, l'association d'informatique de l'ENSTA !



Hébergez vos dépôts Git simplement et en toute sécurité !

A screenshot of a GitLab login form. At the top, there are two tabs: "ENSTA" and "Standard". Below the tabs, there are input fields for "Username or email" and "Password". A "Remember me" checkbox is next to the password field, and a "Forgot your password?" link is to its right. A green "Sign in" button is below the password field. At the bottom, there is a "Sign in with" section with a "cascad ENSTA" button.

End!
Question?